

Using CourseWeb, please turn in your Matlab code and resulting output (enough to show that the program is working) as a PDF document.

1. **Eigenfaces** Chapter 11 of the course notes is on Eigenfaces. For this assignment you can use the files in the directory `old_faces`, which includes some Matlab scripts and a database of 177 face images, each a grayscale .bmp bitmap file of size 64×64 pixels. The face images have been pre-processed so that the background and hair are removed and the faces have similar lighting conditions.

The notes explain how to reshape each face image into a $1 \times 64^2 = 1 \times 4096$ row vector, and collect them into a matrix. The principal components of the matrix then define the main dimensions of variance in the faces. The program `eigenfaces.m` shows how to do this. These principal components are called *eigenfaces*.

The goal of this problem is to apply the same ideas to a new set of 200 faces in the directory `new_faces`. The subdirectory `new_faces/faces` has 200 faces that have been normalized, cropped, and equalized. The subdirectory `new_faces/smiling_faces` has 200 images of the same people, but they are smiling. Each of these images is a grayscale .jpg file with size 193×162 .



Figure 1: The 64×64 face bitmap image `old_faces/face000.bmp`, along with 193×162 jpeg images `new_faces/faces/23a.jpg` and `new_faces/smiling_faces/23b.jpg`. PLEASE USE THE MATLAB `imread()` FUNCTION TO READ IN IMAGES.

Using Matlab, perform the following steps:

- (a) Modify the program `eigenfaces.m` (available in this directory) to use the `new_faces` images instead of the `old_faces` images. Also, modify it to use the Matlab function `imresize` to downsample each of the new faces by a factor of 3, so it is 64×54 (instead of 193×162). Then: PAD both sides of the image with `zeros(64, 5)` so the result is a 64×64 image.
- (b) Create a function that takes as input a string array of filenames of face images, an integer k , and an integer sample size s — and yields the average face and the first k singular values and eigenfaces as output values for a sample of size s .

Apply your function to both the `new_faces/faces` and the `new_faces/smiling_faces` directories, and plot the absolute value of the difference between your average face and (your downsampled version of) the average face provided in the directory. (You can use the `imagesc` function to display images with automatic rescaling of numeric values.)

- (c) If your mean normal face is $\bar{\mathbf{f}}_0$, and your mean smiling face is $\bar{\mathbf{f}}_1$, render (using `imagesc`) the difference $\bar{\mathbf{f}}_0 - \bar{\mathbf{f}}_1$ (the average difference between normal faces and smiling faces).
- (d) Using your downsampled images, perform PCA on each set of faces (normal and smiling). For each image (normal or smiling), construct its $64^2 \times 1$ vector \mathbf{f} . Then, subtract the average face ($\bar{\mathbf{f}}_0$ or $\bar{\mathbf{f}}_1$) and project the remainder on the first $k = 60$ eigenfaces. For example, with a smiling face, the projection of \mathbf{f} on the j -th smiling eigenface \mathbf{e}_j is

$$c_j = \langle (\mathbf{f} - \bar{\mathbf{f}}_1), \mathbf{e}_j \rangle \quad (j = 1, \dots, k).$$

For each set of faces (normal or smiling), make one large scree plot for the set, showing all sequences of the first k coefficients for each image overplotted (e.g. with `hold on`). ????

- (e) Let us say the *unusualness* of a face is the L_2 norm of its eigen-coefficient vector. Determine, for each set, the most unusual face. Normal: Face 124

- (f) Develop three different face classifiers using the eigenfaces you've computed; each should be a function that, given a face image \mathbf{f} as input, yields the output value 1 if \mathbf{f} is smiling, and 0 otherwise. Smiling: 69

Specifically, implement the following 3 classifiers that take an input image \mathbf{f} : N: 1 S: 0

- i. *Classifier X*: yield 1 if the normal face unusualness of \mathbf{f} is less than smiling face unusualness of \mathbf{f} , else 0.
- ii. *Classifier Y*: yield 1 if $\|\mathbf{f} - \mathbf{f}_0\|^2 \geq \|\mathbf{f} - \mathbf{f}_1\|^2$, else 0. N: 1 S: 1
- iii. *Classifier Z*: if C_0 is the covariance matrix for normal faces, and C_1 is the covariance matrix for smiling faces, yield 1 if $\|\mathbf{f} - \mathbf{f}_0\|_{C_0}^2 \geq \|\mathbf{f} - \mathbf{f}_1\|_{C_1}^2$, where $\|\mathbf{x}\|_C^2 = \mathbf{x}' C^{-1} \mathbf{x}$, else 0.

Using each of these three classifiers, determine the classification it yields for the two most unusual images you found in the previous question.

- (g) Write a function `[Sublist1 Sublist2] = randsplit(List)` that takes an array `List` of length n and splits it randomly into two sublists of size `floor(n/2)` and `ceil(n/2)`. (Hint: `randperm`)

Use `randsplit` to split each of the 200-face sets into a training subset and testing subset of equal size.

For both sets of faces (100 normal faces and 100 smiling faces) in the training set, compute: both average faces $\mathbf{f}_0, \mathbf{f}_1$, both face covariance matrices C_0, C_1 , and both sets of eigenfaces.

- (h) For each of the three Classifiers (X, Y, Z) above: error = 0.75, 0.25 x, y
- i. classify each of the 200 faces \mathbf{f} in the testing set, and count classification errors.
 - ii. compute the *error rate* (percentage of errors in test face classifications) for the Classifier.

Then rank the three classifiers by their error rate.

2. Face Compression

In the previous problem you produced a scree plot showing the first 60 eigenface coefficients for each face, and determined the most unusual face. text

For each 64×64 image X from your (downsampled) smiling faces, compute the following sorted sequences:

- (a) sorted absolute values of eigenface coefficients for X
- (b) sorted absolute values of coefficient values in the two-sided FFT of X (in Matlab: `fft2(X)`) DCT 0.0657
- (c) sorted absolute values of coefficient values in the two-sided DCT of X (in Matlab: `dct2(X)`) FFT 0.0448
- (d) sorted singular values from the SVD of X . PCA: 0.0968

We get an *image compression* scheme if we keep only the first k coefficients, and discard the rest.^W

Suppose we keep only the first 60 coefficients. If we compute *compression error* as (the L_2 norm of absolute values of coefficients after the first 60) divided by (the L_2 norm of absolute values of all coefficients), we can compare compression with each of the 4 transforms above.

For the smiling test set (the last 100 smiling images), compute the average compression error for the 4 transforms. Rank the 4 transforms above by their compression error.

3. Floating Point Horror

Using the Matlab command `format long` to make the following floating point results visible.

For each of the following matrices X :

- print the results of `svd(X)`, `log10(svd(X))`, and `range(log10(svd(X)))`, `cond(X)`.
- print the Frobenius norm of the difference between the product $X * \text{inv}(X)$ and the identity matrix.
- Write a program that computes the *exact* inverse X^{-1} of X . (With the `hilb(n)` matrix, this is computed by `invhilb(n)` for example, but you may need to do a little research on the internet to get the form of the exact inverse.) Then print the Frobenius norm of the difference between the product $X * X^{-1}$ and the identity matrix.

- (a) The Hilbert matrix `hilb(12)`.
- (b) The Pascal matrix `pascal(12)`.
- (c) The Fourier matrix `fft(eye(12))`.
- (d) The Vandermonde matrix `vander(1:12)`.